

# 1. Demonstrate use of tensorflow and pytorch by implementing simple code in python

\*\*\*TENSORFLOW\*\*\*

```
import tensorflow as tf

from tensorflow.keras import layers, models

from tensorflow.keras.datasets import mnist

import matplotlib.pyplot as plt

(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Normalize the images to the range of 0 to 1
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Reshape the data to (num_samples, 28 * 28)
x_train = x_train.reshape((x_train.shape[0], 28 * 28))
x_test = x_test.reshape((x_test.shape[0], 28 * 28))

# Create a simple feedforward neural network model
model = models.Sequential([

    layers.Dense(128, activation='relu', input_shape=(28 * 28,)),

    layers.Dense(64, activation='relu'),

    layers.Dense(10, activation='softmax') # Output layer for 10 classes

])

model.compile(optimizer='adam',

              loss='sparse_categorical_crossentropy',

              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.2)

test_loss, test_accuracy = model.evaluate(x_test, y_test)

print(f'Test accuracy (TensorFlow): {test_accuracy:.4f}')

plt.plot(history.history['accuracy'], label='Train Accuracy')

plt.plot(history.history['val_accuracy'], label='Validation Accuracy')

plt.title('Training History (TensorFlow)')

plt.xlabel('Epochs')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend()
```

```
plt.show()
```

```
***Pytorch***
```

```
import torch
```

```
import torch.nn as nn
```

```
import torch.optim as optim
```

```
from torchvision import datasets, transforms
```

```
from torch.utils.data import DataLoader
```

```
import matplotlib.pyplot as plt
```

```
class FeedForwardNN(nn.Module):
```

```
    def __init__(self):
```

```
        super(FeedForwardNN, self).__init__()
```

```
        self.fc1 = nn.Linear(28 * 28, 128)
```

```
        self.fc2 = nn.Linear(128, 64)
```

```
        self.fc3 = nn.Linear(64, 10) # Output layer for 10 classes
```

```
    def forward(self, x):
```

```
        x = x.view(-1, 28 * 28) # Flatten the input
```

```
        x = torch.relu(self.fc1(x))
```

```
        x = torch.relu(self.fc2(x))
```

```
        return self.fc3(x)
```

```
transform = transforms.Compose([transforms.ToTensor()])
```

```
train_dataset = datasets.MNIST(root='./data', train=True, download=True, transform=transform)
```

```
test_dataset = datasets.MNIST(root='./data', train=False, download=True, transform=transform)
```

```
train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True)
```

```
test_loader = DataLoader(dataset=test_dataset, batch_size=32, shuffle=False)
```

```
# Instantiate the model, define the loss function and the optimizer
model = FeedForwardNN()
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters())

# Training settings
num_epochs = 5
train_losses = []

# Train the model
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    for images, labels in train_loader:
        optimizer.zero_grad() # Clear the gradients
        outputs = model(images) # Forward pass
        loss = criterion(outputs, labels) # Compute the loss
        loss.backward() # Backward pass
        optimizer.step() # Update the weights
        running_loss += loss.item()

    # Record the average loss for the epoch
    average_loss = running_loss / len(train_loader)
    train_losses.append(average_loss)
    print(f'Epoch [{epoch + 1}/{num_epochs}], Loss: {average_loss:.4f}')

# Evaluate the model
model.eval()
correct = 0
```

```
total = 0
```

```
with torch.no_grad():
```

```
    for images, labels in test_loader:
```

```
        outputs = model(images)
```

```
        _, predicted = torch.max(outputs.data, 1)
```

```
        total += labels.size(0)
```

```
        correct += (predicted == labels).sum().item()
```

```
print(f'Test accuracy (PyTorch): {correct / total:.4f}')
```

```
# Plot the training loss
```

```
plt.plot(train_losses, label='Training Loss', color='blue')
```

```
plt.title('Training Loss Over Epochs')
```

```
plt.xlabel('Epochs')
```

```
plt.ylabel('Loss')
```

```
plt.legend()
```

```
plt.show()
```

## 2. Implement Feedforward neural networks with Keras and TensorFlow MNIST Digit dataset

*\*\*Importing Dataset from local PC : \*\**

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Load data from CSV file
data = pd.read_csv(r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\mnist_784_csv.csv")

# Separate features and labels
x = data.iloc[:, :-1].values # All columns except the last one are pixel values
y = data['class'].values     # The last column is the class label

# Normalize the pixel values
x = x / 255.0

# Reshape x to (number of samples, 28, 28) to match the input shape for the model
x = x.reshape(-1, 28, 28)

# One-hot encode the labels
y = tf.keras.utils.to_categorical(y, 10)

# Split data into training and testing sets
split_index = int(0.8 * len(x))
x_train, x_test = x[:split_index], x[split_index:]
y_train, y_test = y[:split_index], y[split_index:]

# Build a simple feedforward neural network
model = Sequential([
    Flatten(input_shape=(28, 28)), # Flatten input images to 1D vectors
    Dense(128, activation='relu'),  # Hidden layer with 128 neurons
    Dense(64, activation='relu'),
    Dense(10, activation='softmax') # Output layer with 10 neurons (one for each class)
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=32, validation_data=(x_test, y_test))

# Evaluate the model on test data
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_accuracy:.4f}')
```

```

# Plot training & validation accuracy over epochs
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Display a few random test images with their predictions
num_images = 5
random_indices = np.random.choice(x_test.shape[0], num_images, replace=False)

for i, idx in enumerate(random_indices):
    test_image = x_test[idx]
    true_label = np.argmax(y_test[idx])

    # Predict the label for the test image
    test_image_resaped = np.expand_dims(test_image, axis=0) # Reshape for prediction
    predicted_label = np.argmax(model.predict(test_image_resaped))

    # Plot the image and prediction
    plt.subplot(1, num_images, i + 1)
    plt.imshow(test_image, cmap='gray')
    plt.title(f"True: {true_label}\nPred: {predicted_label}")
    plt.axis('off')
plt.show()

```

.....

*\*\*Importing Dataset Internet : \*\**

```

import tensorflow as tf
import numpy as np
from tensorflow.keras import layers, models
import tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
y_train, y_test = to_categorical(y_train, 10), to_categorical(y_test, 10)

model = Sequential([
    Flatten(input_shape = (28, 28)),
    Dense(256, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(64, activation = 'relu'),
    Dense(10, activation = 'softmax')
])

```

```
)  
model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])  
  
history = model.fit(x_train, y_train, epochs = 10, batch_size=32, validation_data = (x_test, y_test))  
test_loss, test_accuracy = model.evaluate(x_test, y_test)  
  
print(f'Test accuracy: {test_accuracy: .4f}')  
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model Accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend(['Train', 'Test'], loc = 'upper left' )  
plt.show()  
  
random_index = np.random.randint(0, x_test.shape[0])  
test_image = x_test[random_index]  
true_label = np.argmax(y_test[random_index])  
test_image_reshaped = np.expand_dims(test_image, axis = 0)  
predicted_probabilities = model.predict(test_image_reshaped)  
predicted_label = np.argmax(predicted_probabilities)  
  
plt.imshow(test_image, cmap = 'gray')  
plt.title(f"True Label: {true_label}, Predicted Label: {predicted_label}")  
plt.axis('off')  
plt.show()
```

.....





### 3. Implement Feedforward neural networks with Keras and TensorFlow CIFAR dataset

```
import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Flatten

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

from sklearn.metrics import confusion_matrix, classification_report

import seaborn as sns


class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

train_data = pd.read_csv('path_to_your_training_file.csv')

x_train = train_data.iloc[:, :-1].values # All columns except the last one are pixel values

y_train = train_data['label'].values # The last column is the label

test_data = pd.read_csv('path_to_your_testing_file.csv')

x_test = test_data.iloc[:, :-1].values

y_test = test_data['label'].values


# Normalize the pixel values

x_train = x_train / 255.0

x_test = x_test / 255.0


# Reshape the data to match CIFAR-10 image dimensions (32, 32, 3)

x_train = x_train.reshape(-1, 32, 32, 3)

x_test = x_test.reshape(-1, 32, 32, 3)


# One-hot encode the labels

y_train_encoded = tf.keras.utils.to_categorical(y_train, 10)

y_test_encoded = tf.keras.utils.to_categorical(y_test, 10)
```

```

model = Sequential([
    Flatten(input_shape=(32, 32, 3)), # Flatten 32x32x3 images
    Dense(512, activation='relu'),
    Dense(256, activation='relu'),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax') # 10 output classes for CIFAR-10
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(x_train, y_train_encoded, epochs=10, batch_size=32, validation_split=0.2)

test_loss, test_accuracy = model.evaluate(x_test, y_test_encoded)

print(f'Test accuracy: {test_accuracy:.4f}')

plt.figure(figsize=(12, 4))

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Function to visualize a random prediction
def visualize_random_prediction(images, true_labels, model):
    random_index = np.random.randint(0, len(images))
    image = images[random_index]
    true_label = true_labels[random_index]

```

```
# Prepare the image for prediction
```

```
image_for_prediction = np.expand_dims(image, axis=0)
```

```
prediction = model.predict(image_for_prediction)
```

```
# Get class names for true and predicted labels
```

```
true_class = class_names[np.argmax(true_label)]
```

```
predicted_class = class_names[np.argmax(prediction)]
```

```
# Plot the image and prediction
```

```
plt.figure(figsize=(4, 4))
```

```
plt.imshow(image)
```

```
plt.title(f"True Label: {true_class}\nPredicted Label: {predicted_class}",
```

```
        color='green' if true_class == predicted_class else 'red')
```

```
plt.axis('off')
```

```
plt.show()
```

```
# Visualize a random test image and its prediction
```

```
visualize_random_prediction(x_test, y_test_encoded, model)
```

```
# Function to plot confusion matrix and classification report
```

```
def plot_confusion_matrix(y_true, y_pred):
```

```
    cm = confusion_matrix(y_true, y_pred)
```

```
    plt.figure(figsize=(10, 8))
```

```
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
```

```
                xticklabels=class_names,
```

```
                yticklabels=class_names)
```

```
    plt.title('Confusion Matrix')
```

```
    plt.xlabel('Predicted Label')
```

```
    plt.ylabel('True Label')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# Predict on the test set and display confusion matrix
```

```
predictions = model.predict(x_test)
```

```
y_pred_classes = np.argmax(predictions, axis=1)
```

```
y_true_classes = y_test # Use original test labels for evaluation
```

```
plot_confusion_matrix(y_true_classes, y_pred_classes)
```

```
# Print classification report
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_true_classes, y_pred_classes, target_names=class_names))
```

#### 4. Build image classification model using CNN on fashion MNIST dataset.

```
import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import random
from sklearn.metrics import classification_report, confusion_matrix
from tensorflow.keras.optimizers import Adam

train_data = pd.read_csv(r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\fashion-
mnist_train.csv\fashion-mnist_train.csv")
test_data = pd.read_csv(r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\fashion-
mnist_test.csv\fashion-mnist_test.csv")

# Prepare the data
X_train = train_data.iloc[:, 1:].values # Get pixel values
y_train = train_data.iloc[:, 0].values # Get labels
X_test = test_data.iloc[:, 1:].values # Get pixel values
y_test = test_data.iloc[:, 0].values # Get labels

# Normalize the data
X_train = X_train / 255.0
X_test = X_test / 255.0

# Reshape the data to match the model input
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)

# Define the model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dense(10, activation='softmax')
])

model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```

# Train the model
history = model.fit(X_train, y_train, epochs=30, validation_split=0.2)

# Plot training history
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='upper left')
plt.show()

# Evaluate the model on the test data
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f'Test accuracy: {test_accuracy:.4f}')

# Predictions and classification report
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred_classes))

def show_random_prediction(X_test, y_test, model):
    random_index = random.randint(0, len(X_test) - 1)
    image = X_test[random_index]
    true_label = y_test[random_index]
    image_reshaped = np.expand_dims(image, axis=0)

    # Predict the class of the image
    prediction = model.predict(image_reshaped)
    predicted_label = np.argmax(prediction)

    plt.imshow(image.squeeze(), cmap='gray')
    plt.title(f"True Label: {true_label}\nPredicted Label: {predicted_label}")
    plt.axis('off')
    plt.show()

# Call the function to display a random prediction
show_random_prediction(X_test, y_test, model)

```

5. Build image classification model using CNN on pneumonia X RAY IMAGE dataset.

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import numpy as np
import os

train_dir = ""
val_dir = ""
test_dir = ""

# Image data generators for loading images in batches
train_datagen = ImageDataGenerator(rescale=1./255)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Generate batches of image data from directories
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150), # Resize images to 150x150 pixels
    batch_size=32,
    class_mode='binary' # Binary classification for 'Normal' and 'Pneumonia'
)

val_generator = val_datagen.flow_from_directory(
    val_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary'
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary'
)

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
```

```

layers.Flatten(),
layers.Dense(512, activation='relu'),
layers.Dense(1, activation='sigmoid') # Sigmoid for binary classification
])

```

```

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

```

```

history = model.fit(
    train_generator,
    steps_per_epoch=len(train_generator),
    epochs=10,
    validation_data=val_generator,
    validation_steps=len(val_generator)
)

```

```

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.title('Model Accuracy')
plt.tight_layout()
plt.show()

```

```

test_loss, test_accuracy = model.evaluate(test_generator)
print(f'Test accuracy: {test_accuracy:.4f}')

```

```

from tensorflow.keras.preprocessing import image

```

```

def predict_image(img_path, model):
    img = image.load_img(img_path, target_size=(150, 150))
    img_array = image.img_to_array(img) / 255.0
    img_array = np.expand_dims(img_array, axis=0) # Add batch dimension

```

```

    prediction = model.predict(img_array)
    class_label = 'Pneumonia' if prediction[0] > 0.5 else 'Normal'
    print(f"Prediction: {class_label}")

```

```

    plt.imshow(img)
    plt.title(f"Predicted Label: {class_label}")
    plt.axis('off')
    plt.show()

```

```

# Call the function with the path to a test image
predict_image("")

```



6. Build image classification model using CNN on FOOD dataset.

7. Build Brain tumor classification model with CNN

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
import matplotlib.pyplot as plt
```

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
import numpy as np
```

```
from tensorflow.keras.preprocessing import image
```

```
import matplotlib.pyplot as plt
```

```
train_dir = r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\brain tumour\Training"
```

```
test_dir = r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\brain tumour\Testing"
```

```
train_datagen = ImageDataGenerator(
```

```
    rescale=1.0/255,
```

```
    rotation_range=30,
```

```
    width_shift_range=0.2,
```

```
    height_shift_range=0.2,
```

```
    shear_range=0.2,
```

```
    zoom_range=0.2,
```

```
    horizontal_flip=True,
```

```
    vertical_flip=True,
```

```
    fill_mode='nearest',
```

```
    validation_split=0.2
```

```
)
```

```
test_datagen = ImageDataGenerator(rescale=1.0/255)
```

```
train_generator = train_datagen.flow_from_directory(
```

```
    train_dir,
```

```
    target_size=(150, 150),
```

```

    batch_size=32,
    class_mode='sparse',
    subset='training'
)

val_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='sparse',
    subset='validation'
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='sparse'
)

model = models.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3)),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3,3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(128, (3,3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(4, activation='softmax')
])

```

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

history = model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=30,
    callbacks=[early_stopping]
)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plot loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

```

plt.show()

import numpy as np

from tensorflow.keras.preprocessing import image

import matplotlib.pyplot as plt

# Define a function to predict a single image
def predict_single_image(model, img_path, class_labels):
    # Load the image with target size as expected by the model
    img = image.load_img(img_path, target_size=(150, 150))

    # Convert the image to an array and rescale it
    img_array = image.img_to_array(img) / 255.0
    img_array = np.expand_dims(img_array, axis=0) # Expand dimensions for batch shape

    # Make a prediction
    prediction = model.predict(img_array)
    predicted_class = class_labels[np.argmax(prediction)]
    confidence = np.max(prediction)

    # Display the image with prediction
    plt.imshow(img)
    plt.title(f"Predicted: {predicted_class} ({confidence:.2f})")
    plt.axis('off')
    plt.show()

# Class labels in order (change these if necessary)
class_labels = ['no_tumor', 'pituitary_tumor', 'glioma', 'meningioma']

#Example
predict_single_image(model, r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\brain
tumour\Testing\pituitary_tumor\image(45).jpg", class_labels)

```

## 8. Build Recurrent Neural Network by using the numpy library

```
import numpy as np
import matplotlib.pyplot as plt
```

### Dataset Creation

```
def create_dataset():
    data = "hello world hello world hello world"
    chars = list(set(data))
    char_to_idx = {ch: i for i, ch in enumerate(chars)}
    idx_to_char = {i: ch for i, ch in enumerate(chars)}

    # Prepare training data
    X = [] # Input sequences
    y = [] # Target sequences

    for i in range(len(data) - 1):
        input_char = np.zeros((len(chars)))
        input_char[char_to_idx[data[i]]] = 1
        target_char = np.zeros((len(chars)))
        target_char[char_to_idx[data[i + 1]]] = 1

        X.append(input_char)
        y.append(target_char)

    return np.array(X), np.array(y), char_to_idx, idx_to_char, data
```

### Model Parameters

```
class SimpleRNN:
    def __init__(self, input_size, hidden_size, output_size, learning_rate=0.01):
        # Initialize weights
        self.Wxh = np.random.randn(hidden_size, input_size) * 0.01 # Input to hidden
        self.Whh = np.random.randn(hidden_size, hidden_size) * 0.01 # Hidden to hidden
        self.Why = np.random.randn(output_size, hidden_size) * 0.01 # Hidden to output

        # Initialize biases
        self.bh = np.zeros((hidden_size, 1)) # Hidden bias
        self.by = np.zeros((output_size, 1)) # Output bias

        self.learning_rate = learning_rate

    def sigmoid(self, x):
        return 1 / (1 + np.exp(-x))

    def tanh(self, x):
        return np.tanh(x)

    def tanh_derivative(self, x):
        return 1 - np.tanh(x)**2
```

```

def forward(self, inputs):
    # Initialize lists to store states
    self.hidden_states = []
    self.outputs = []
    h_prev = np.zeros((self.Whh.shape[0], 1)) # Initial hidden state

    # Forward pass for each time step
    for x in inputs:
        # Convert input to column vector
        x = x.reshape(-1, 1)

        # Calculate hidden state
        h = self.tanh(np.dot(self.Wxh, x) + np.dot(self.Whh, h_prev) + self.bh)

        # Calculate output
        y = self.sigmoid(np.dot(self.Why, h) + self.by)

        # Store states for backpropagation
        self.hidden_states.append(h)
        self.outputs.append(y)
        h_prev = h

    return self.outputs

```

```

def backward(self, inputs, targets, outputs, hidden_states):
    # Initialize gradients
    dWxh = np.zeros_like(self.Wxh)
    dWhh = np.zeros_like(self.Whh)
    dWhy = np.zeros_like(self.Why)
    dbh = np.zeros_like(self.bh)
    dby = np.zeros_like(self.by)

    dh_next = np.zeros_like(hidden_states[0])

    # Backpropagate through time
    for t in reversed(range(len(outputs))):
        dy = outputs[t] - targets[t].reshape(-1, 1)

        # Gradients for Why and by
        dWhy += np.dot(dy, hidden_states[t].T)
        dby += dy

        # Gradient for hidden state
        dh = np.dot(self.Why.T, dy) + dh_next

        # Gradient through tanh
        dh_raw = self.tanh_derivative(hidden_states[t]) * dh

        dbh += dh_raw
        dWxh += np.dot(dh_raw, inputs[t].reshape(1, -1))
        dWhh += np.dot(dh_raw, hidden_states[t-1].T) if t > 0 else np.dot(dh_raw,
np.zeros_like(hidden_states[0]).T)

```

```

dh_next = np.dot(self.Whh.T, dh_raw)

# Clip gradients to prevent exploding gradients
for dparam in [dWxh, dWhh, dWhy, dbh, dby]:
    np.clip(dparam, -5, 5, out=dparam)

# Update weights and biases
self.Wxh -= self.learning_rate * dWxh
self.Whh -= self.learning_rate * dWhh
self.Why -= self.learning_rate * dWhy
self.bh -= self.learning_rate * dbh
self.by -= self.learning_rate * dby

```

### Training and History

```

def train_and_demonstrate():
    # Create dataset
    X, y, char_to_idx, idx_to_char, original_data = create_dataset()

    # Initialize RNN
    input_size = len(char_to_idx)
    hidden_size = 50
    output_size = len(char_to_idx)
    rnn = SimpleRNN(input_size, hidden_size, output_size)

    # Training loop
    epochs = 100
    losses = []

    print("Training the RNN...")
    print("Original sequence:", original_data)
    print("\nTraining Progress:")

    for epoch in range(epochs):
        # Forward pass
        outputs = rnn.forward(X)

        # Calculate loss (mean squared error)
        loss = np.mean([(output - target.reshape(-1, 1))**2 for output, target in zip(outputs, y)])
        losses.append(loss)

        # Backward pass
        rnn.backward(X, y, outputs, rnn.hidden_states)

        if epoch % 20 == 0 or epoch == epochs - 1:
            print(f'Epoch {epoch}, Loss: {loss:.4f}')

```

### Demonstration and Visualization

```

print("\nDemonstrating predictions for each character in sequence:")
for i in range(len(original_data) - 1):
    input_char = original_data[i]
    actual_next_char = original_data[i + 1]

```

```

# Prepare input
input_vector = np.zeros((len(char_to_idx)))
input_vector[char_to_idx[input_char]] = 1

# Get prediction
output = rnn.forward([input_vector])[0]
predicted_char = idx_to_char[np.argmax(output)]

print(f"Input: '{input_char}' → Predicted next: '{predicted_char}' (Actual: '{actual_next_char}')"

# Plot training loss
plt.figure(figsize=(10, 5))
plt.plot(losses)
plt.title('Training Loss over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True)
plt.show()

return rnn, char_to_idx, idx_to_char

```

#### Run the Demonstration

```
rnn, char_to_idx, idx_to_char = train_and_demonstrate()
```



9. Implement simple autoencoder to reconstruct MNIST digits. Add sparsity constraint on the encoded representations

```
import numpy as np
import pandas as pd # Import pandas for data loading
import tensorflow as tf
from tensorflow.keras import layers, regularizers, models
import matplotlib.pyplot as plt

data = pd.read_csv(r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\mnist_784_csv.csv")
x_data = data.iloc[:, 1:].values # Assuming the first column is labels and the rest are pixel values

# Normalize the pixel values
x_data = x_data.astype('float32') / 255.

# Assuming each image is 28x28
num_samples = x_data.shape[0] # Get the number of samples
x_data = np.reshape(x_data, (num_samples, 28, 28, 1)) # Reshape to (num_samples, 28, 28, 1)

# Split into training and testing datasets
x_train, x_test = x_data[:int(len(x_data) * 0.8)], x_data[int(len(x_data) * 0.8):]

# Parameters
input_shape = (28, 28, 1) # Change this according to your data's shape
encoding_dim = 128
sparsity_penalty = 1e-5 # L1 sparsity regularization parameter

# Define encoder
def build_encoder(input_shape, encoding_dim, sparsity_penalty):
    inputs = layers.Input(shape=input_shape)
    x = layers.Flatten()(inputs)
    x = layers.Dense(256, activation='relu')(x)
    x = layers.BatchNormalization()(x)

    # Apply L1 regularization to encourage sparsity
    encoded = layers.Dense(encoding_dim,
                           activation='relu',
                           activity_regularizer=regularizers.L1(sparsity_penalty))(x)
    return models.Model(inputs, encoded, name="encoder")

# Define decoder
def build_decoder(encoding_dim, original_shape):
    encoded_input = layers.Input(shape=(encoding_dim,))
    x = layers.Dense(256, activation='relu')(encoded_input)
    x = layers.BatchNormalization()(x)
    x = layers.Dense(np.prod(original_shape), activation='sigmoid')(x)
    decoded = layers.Reshape(original_shape)(x)

    return models.Model(encoded_input, decoded, name="decoder")

# Build autoencoder
encoder = build_encoder(input_shape, encoding_dim, sparsity_penalty)
```

```

decoder = build_decoder(encoding_dim, input_shape)
autoencoder = models.Model(encoder.input, decoder(encoder.output), name="autoencoder")

# Compile model
autoencoder.compile(optimizer='adam',
                    loss='binary_crossentropy')

# Train autoencoder
history = autoencoder.fit(x_train, x_train,
                        epochs=30,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test, x_test))

# Generate reconstructions
decoded_imgs = autoencoder.predict(x_test)

# Visualize results
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title('Original')
    plt.axis('off')

    # Display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    plt.title('Reconstructed')
    plt.axis('off')

plt.show()

# Plot training history
plt.figure(figsize=(10, 4))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

## 10. Use Autoencoder to implement anomaly detection on credit card dataset

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt

data = pd.read_csv(r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\creditcard.csv")

# Separate features and labels
X = data.drop('Class', axis=1)
y = data['Class']

# Normalize the feature data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Filter normal transactions (Class == 0) for training
X_train_normal = X_train[y_train == 0]
print(f"Number of normal transactions in training set: {X_train_normal.shape[0]}")

input_dim = X_train_normal.shape[1]

autoencoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(8, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='sigmoid')
])

autoencoder.compile(optimizer='adam', loss='mse')

# Train the autoencoder on normal transactions
history = autoencoder.fit(X_train_normal, X_train_normal,
                        epochs=50,
                        batch_size=256,
                        validation_split=0.2,
                        shuffle=True)

# Calculate reconstruction error on the test set
X_test_predictions = autoencoder.predict(X_test)
mse = np.mean(np.square(X_test - X_test_predictions), axis=1)
```

```

# Set a threshold for anomaly detection (95th percentile of MSE on normal transactions)
threshold = np.percentile(mse[y_test == 0], 95)
print(f"Threshold for anomaly detection: {threshold}")

# Classify anomalies
y_pred = (mse > threshold).astype(int)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Visualization of reconstruction error distribution
plt.figure(figsize=(10, 6))
plt.hist(mse[y_test == 0], bins=50, alpha=0.6, label='Normal')
plt.hist(mse[y_test == 1], bins=50, alpha=0.6, label='Fraud')
plt.axvline(threshold, color='r', linestyle='--', label='Threshold')
plt.title("Reconstruction Error Distribution")
plt.xlabel("Reconstruction error")
plt.ylabel("Frequency")
plt.legend()
plt.show()

# Show reconstruction error and prediction for a few samples
sample_indices = [0, 1, 2, 3, 4] # Choose some random test indices for demonstration
for index in sample_indices:
    sample = X_test[index]
    reconstruction = autoencoder.predict(sample.reshape(1, -1))
    error = np.mean(np.square(sample - reconstruction))
    prediction = "Fraud" if error > threshold else "Normal"
    print(f"Sample {index}:")
    print(f" Reconstruction Error: {error}")
    print(f" Prediction: {prediction} (Threshold: {threshold})")

```

## #OUTPUT

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.95	0.97	56864
1	0.03	0.90	0.06	98
accuracy			0.95	56962
macro avg	0.51	0.92	0.52	56962
weighted avg	1.00	0.95	0.97	56962

- **Precision:**

- For class 0 (normal transactions), precision is 1.00, meaning that when the model predicts a transaction is normal, it is correct 100% of the time.
- For class 1 (fraudulent transactions), precision is only 0.03. This means that of all the transactions predicted as fraudulent, only 3% were actually fraudulent. This indicates a high number of false positives.

- **Recall:**

- For class 0, recall is 0.95, meaning the model correctly identifies 95% of the actual normal transactions.
- For class 1, recall is 0.90, indicating that 90% of the fraudulent transactions are correctly identified by the model. This shows that the model is effective at identifying actual fraud cases, despite the low precision.

- **F1-Score:** This is the harmonic mean of precision and recall.

- The F1-score for normal transactions is quite high (0.97), indicating that the model performs very well in detecting normal transactions.
- The F1-score for fraudulent transactions is low (0.06), highlighting the difficulty in identifying fraud without a significant number of false positives.

Confusion Matrix:

```
[[54020 2844]
 [ 10  88]]
```

□ The first row corresponds to normal transactions (Class 0):

- **54020** true negatives (correctly predicted as normal).
- **2844** false positives (incorrectly predicted as fraud).

□ The second row corresponds to fraudulent transactions (Class 1):

- **10** false negatives (incorrectly predicted as normal).
- **88** true positives (correctly predicted as fraud).

1/1 [=====] - 0s 22ms/step

Sample 0:

Reconstruction Error: 65.5436218925837

Prediction: Fraud (Threshold: 1.4906384262057126)

1/1 [=====] - 0s 19ms/step

Sample 1:

Reconstruction Error: 0.8376235550297751

Prediction: Normal (Threshold: 1.4906384262057126)

1/1 [=====] - 0s 30ms/step

Sample 2:

Reconstruction Error: 0.31452216628550395

Prediction: Normal (Threshold: 1.4906384262057126)

1/1 [=====] - 0s 18ms/step

Sample 3:

Reconstruction Error: 0.2853646274277654

Prediction: Normal (Threshold: 1.4906384262057126)



## 11. Implement the concept of image denoising using autoencoders on MNIST data set

\*\*\**Import Data from Local PC*\*\*\*

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.layers import Input, Conv2D, Conv2DTranspose
import matplotlib.pyplot as plt
import pandas as pd

data = pd.read_csv('your_file_path.csv')

# Extract features and labels
X = data.iloc[:, :-1].values.reshape(-1, 28, 28, 1)
y = data.iloc[:, -1].values

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# Add noise to the images
noise_factor = 0.5
X_train_noisy = X_train + noise_factor * np.random.normal(loc=0.0, scale=1.0,
size=X_train.shape)
X_test_noisy = X_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=X_test.shape)
X_train_noisy = np.clip(X_train_noisy, 0.0, 1.0)
X_test_noisy = np.clip(X_test_noisy, 0.0, 1.0)

# Define the autoencoder model
input_img = Input(shape=(28, 28, 1))
x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
encoded = Conv2D(64, (3, 3), activation='relu', padding='same')(x)

x = Conv2DTranspose(64, (3, 3), activation='relu', padding='same')(encoded)
x = Conv2DTranspose(32, (3, 3), activation='relu', padding='same')(x)
x = Conv2DTranspose(32, (3, 3), activation='relu', padding='same')(x)
decoded = Conv2DTranspose(1, (3, 3), activation='sigmoid', padding='same')(x)

autoencoder = models.Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.summary()

# Train the autoencoder
autoencoder.fit(X_train_noisy, X_train, epochs=10, batch_size=128, shuffle=True,
validation_data=(X_test_noisy, X_test))
```

.....

```

# Predict denoised images
denoised_images = autoencoder.predict(X_test_noisy)

# Visualize results
n = 10
plt.figure(figsize=(20, 6))
for i in range(n):
    # Display Original
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(X_test[i].reshape(28, 28), cmap='gray')
    plt.title("Original")
    plt.axis("off")

    # Display Noisy
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(X_test_noisy[i].reshape(28, 28), cmap='gray')
    plt.title("Noisy")
    plt.axis("off")

    # Display Denoised
    ax = plt.subplot(3, n, i + 1 + 2 * n)
    plt.imshow(denoised_images[i].reshape(28, 28), cmap='gray')
    plt.title("Denoised")
    plt.axis("off")

plt.show()

```

.....

### \*\*\*Import from Internet\*\*\*

```

import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.layers import Input, Conv2D, Conv2DTranspose
import matplotlib.pyplot as plt

(x_train, _), (x_test, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

x_train = np.reshape(x_train, (len(x_train), 28, 28, 1))
x_test = np.reshape(x_test, (len(x_test), 28, 28, 1))

noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)
x_train_noisy = np.clip(x_train_noisy, 0.0, 1.0)
x_test_noisy = np.clip(x_test_noisy, 0.0, 1.0)

input_img = Input(shape=(28, 28, 1))
x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)

```



```
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x) # Removed strides here
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x) # Removed strides here
encoded = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
```

```
# Decoder
```

```
x = Conv2DTranspose(64, (3, 3), activation='relu', padding='same')(encoded)
x = Conv2DTranspose(32, (3, 3), activation='relu', padding='same')(x)
x = Conv2DTranspose(32, (3, 3), activation='relu', padding='same')(x)
decoded = Conv2DTranspose(1, (3, 3), activation='sigmoid', padding='same')(x)
```

```
autoencoder = models.Model(input_img, decoded)
```

```
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.summary()
```

```
autoencoder.fit(x_train_noisy, x_train, epochs=30, batch_size=128, shuffle=True,
validation_data=(x_test_noisy, x_test))
```

```
denoised_images = autoencoder.predict(x_test_noisy)
```

```
n = 10
```

```
plt.figure(figsize=(20, 6))
```

```
for i in range(n):
```

```
    # Display Original
```

```
    ax = plt.subplot(3, n, i + 1)
```

```
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
```

```
    plt.title("Original")
```

```
    plt.axis("off")
```

```
    # Display Noisy
```

```
    ax = plt.subplot(3, n, i + 1 + n)
```

```
    plt.imshow(x_test_noisy[i].reshape(28, 28), cmap='gray')
```

```
    plt.title("Noisy")
```

```
    plt.axis("off")
```

```
    # Display Denoised
```

```
    ax = plt.subplot(3, n, i + 1 + 2 * n)
```

```
    plt.imshow(denoised_images[i].reshape(28, 28), cmap='gray')
```

```
    plt.title("Denoised")
```

```
    plt.axis("off")
```

```
plt.show()
```



## 12. Implement object detection using Transfer learning on food dataset

```
import tensorflow as tf
from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras.applications import VGG16
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

train_dir = r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\animal10\train_dataset"
val_dir = r"C:\Users\AVNISH\Desktop\Sem 7\z_DL_Dataset\animal10\val_dataset"
img_size = (224, 224)
batch_size = 16

train_dataset = image_dataset_from_directory(
    train_dir,
    image_size=img_size,
    batch_size=batch_size
)

val_dataset = image_dataset_from_directory(
    val_dir,
    image_size=img_size,
    batch_size=batch_size
)

base_model = VGG16(input_shape=img_size + (3,), include_top=False, weights='imagenet')
base_model.trainable = False

model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.3),
    layers.Dense(len(train_dataset.class_names), activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

history = model.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=10,
    callbacks=[early_stopping]
)

loss, accuracy = model.evaluate(val_dataset)
print(f"Validation accuracy: {accuracy * 100:.2f}%")
```

```
y_true = np.concatenate([y for x, y in val_dataset], axis=0)
y_pred = np.argmax(model.predict(val_dataset), axis=1)
cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(10, 8))
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=val_dataset.class_names)
disp.plot(cmap='Blues', values_format='d')
plt.title("Confusion Matrix")
plt.show()

plt.figure(figsize=(8, 6))
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['Training Accuracy', 'Validation Accuracy'])
plt.title("Training and Validation Accuracy")
plt.show()
```

### 13. Implement the Continuous Bag of Words (CBOW) Model.

```
import numpy as np, import tensorflow as tf, import matplotlib.pyplot as plt
from collections import Counter
from sklearn.decomposition import PCA

# Larger sample text to train on
text = """
Knowledge is the pathway to discovery. Learning transforms our understanding of the world...
"""

# Tokenize and preprocess the text
words = text.lower().split()
vocab = set(words)
vocab_size = len(vocab)

# Create word to index and index to word dictionaries
word_to_ix = {word: i for i, word in enumerate(vocab)}
ix_to_word = {i: word for word, i in word_to_ix.items()}

# Define context window size
context_size = 2 # Larger context window would provide richer context

# Generate context-target pairs
def generate_context_target_pairs(words, context_size):
    pairs = []
    for i in range(context_size, len(words) - context_size):
        context = words[i - context_size:i] + words[i + 1:i + context_size + 1]
        target = words[i]
        pairs.append((context, target))
    return pairs

# Create context-target pairs
data = generate_context_target_pairs(words, context_size)

# Prepare the training data
X = []
y = []

for context, target in data:
    X.append([word_to_ix[w] for w in context])
    y.append(word_to_ix[target])

X = np.array(X)
y = np.array(y)

# CBOW Model Definition
class CBOW(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim):
        super(CBOW, self).__init__()
        self.embeddings = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.linear = tf.keras.layers.Dense(vocab_size)
```

```

def call(self, context_words):
    embedded = self.embeddings(context_words)
    context_vector = tf.reduce_mean(embedded, axis=1) # Average the embeddings of context
words
    out = self.linear(context_vector)
    return out

embedding_dim = 50 # Increase embedding dimension for more nuanced representations
model = CBOW(vocab_size, embedding_dim)
loss_function = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)

epochs = 100
for epoch in range(epochs):
    with tf.GradientTape() as tape:
        predictions = model(X)
        loss = loss_function(y, predictions)

    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))

    if epoch % 20 == 0:
        print(f"Epoch: {epoch}, Loss: {loss.numpy():.4f}")

embeddings = model.embeddings.weights[0].numpy()

# Function to find similar words based on cosine similarity
def find_similar_words(word, embeddings, word_to_ix, ix_to_word, top_n=5):
    word_vec = embeddings[word_to_ix[word]]
    similarities = np.dot(embeddings, word_vec) / (
        np.linalg.norm(embeddings, axis=1) * np.linalg.norm(word_vec) + 1e-9
    )
    closest_words = similarities.argsort()[-(top_n + 1):-1] # Exclude the word itself
    return [ix_to_word[idx] for idx in closest_words]

test_word = "creativity"
similar_words = find_similar_words(test_word, embeddings, word_to_ix, ix_to_word)
print(f"Words similar to '{test_word}': {similar_words}")

def visualize_embeddings(embeddings, word_to_ix):
    pca = PCA(n_components=2)
    reduced_embeddings = pca.fit_transform(embeddings)
    plt.figure(figsize=(10, 10))
    for word, idx in word_to_ix.items():
        x, y = reduced_embeddings[idx]
        plt.scatter(x, y)
        plt.text(x, y, word, fontsize=12)
    plt.show()

visualize_embeddings(embeddings, word_to_ix)

```

## 14. Implement object detection using Transfer learning on food dataset

```
pip install tensorflow
import tensorflow as tf
from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np

train_dir = "path/to/train"
val_dir = "path/to/val"
test_dir = "path/to/test"

# Define parameters
img_size = (224, 224) # MobileNetV2 expects 224x224 images
batch_size = 32

# Load datasets
train_dataset = image_dataset_from_directory(train_dir, image_size=img_size, batch_size=batch_size)
val_dataset = image_dataset_from_directory(val_dir, image_size=img_size, batch_size=batch_size)
test_dataset = image_dataset_from_directory(test_dir, image_size=img_size, batch_size=batch_size)

# Load the base model (MobileNetV2) with pre-trained weights
base_model = tf.keras.applications.MobileNetV2(input_shape=(224, 224, 3),
                                                include_top=False,
                                                weights='imagenet')
base_model.trainable = False # Freeze the base model layers

# Build the model
model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(1, activation='sigmoid') # Binary classification output layer
])

# Compile the model
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.summary()
epochs = 10
history = model.fit(train_dataset, validation_data=val_dataset, epochs=epochs)
test_loss, test_accuracy = model.evaluate(test_dataset)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
print(f"Test Loss: {test_loss:.4f}")
import cv2
import matplotlib.pyplot as plt

def predict_and_display(model, img_path):
    img = tf.keras.preprocessing.image.load_img(img_path, target_size=img_size)
    img_array = tf.keras.preprocessing.image.img_to_array(img)
```

```

img_array = tf.expand_dims(img_array, 0) # Create a batch

prediction = model.predict(img_array)
class_name = 'Dog' if prediction[0] > 0.5 else 'Cat'

plt.imshow(img)
plt.title(f'Prediction: {class_name}')
plt.axis('off')
plt.show()

# Test the function with an example image
predict_and_display(model, "path/to/some_test_image.jpg")
# Plot training & validation accuracy values
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train')
plt.plot(history.history['val_accuracy'], label='Validation')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train')
plt.plot(history.history['val_loss'], label='Validation')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(loc='upper right')

plt.show()
***FOOD DATASET***
import tensorflow as tf
import numpy as np
from object_detection.utils import config_util
from object_detection.builders import model_builder
from object_detection.utils import visualization_utils as viz_utils
import cv2
import matplotlib.pyplot as plt
# Define paths
pipeline_config_path = 'path/to/your/pipeline.config' # Customize with your pipeline config path
model_checkpoint = 'path/to/your/pre-trained-model/checkpoint' # Model checkpoint folder

# Load the pipeline config and build the model
configs = config_util.get_configs_from_pipeline_file(pipeline_config_path)
model_config = configs['model']
detection_model = model_builder.build(model_config=model_config, is_training=True)
# Restore checkpoint for transfer learning
ckpt = tf.compat.v2.train.Checkpoint(model=detection_model)
ckpt.restore(model_checkpoint).expect_partial()

```



```

# Load dataset (example with TFRecord files, adjust accordingly if using COCO format)
train_dataset = tf.data.TFRecordDataset("path/to/train.record")
val_dataset = tf.data.TFRecordDataset("path/to/val.record")

# Parse the dataset using a parser function
def parse_tfrecord_fn(example):
    # Define the parsing schema as per your dataset
    pass # Complete this based on your TFRecord schema

train_dataset = train_dataset.map(parse_tfrecord_fn)
val_dataset = val_dataset.map(parse_tfrecord_fn)
# Define optimizer and metrics
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)

# Training loop
epochs = 10
for epoch in range(epochs):
    print(f'Epoch {epoch + 1}/{epochs}')

    for batch in train_dataset:
        with tf.GradientTape() as tape:
            predictions = detection_model(batch['image'], training=True)
            loss_value = detection_model.compute_loss(predictions, batch['groundtruth_boxes'],
batch['groundtruth_classes'])

            gradients = tape.gradient(loss_value, detection_model.trainable_variables)
            optimizer.apply_gradients(zip(gradients, detection_model.trainable_variables))

        print(f'Loss: {loss_value.numpy()}')

    # Validate on validation dataset if desired
def load_image_into_numpy_array(path):
    return np.array(cv2.imread(path))

def detect_objects_in_image(image_path):
    image_np = load_image_into_numpy_array(image_path)
    input_tensor = tf.convert_to_tensor(image_np)
    input_tensor = input_tensor[tf.newaxis, ...]

    detections = detection_model(input_tensor, training=False)

# Visualize results
viz_utils.visualize_boxes_and_labels_on_image_array(
    image_np,
    detections['detection_boxes'][0].numpy(),
    detections['detection_classes'][0].numpy().astype(np.int32),
    detections['detection_scores'][0].numpy(),
    category_index, # Dictionary mapping class IDs to names (e.g., {1: "apple", 2: "banana"})
    use_normalized_coordinates=True,
    line_thickness=8
)

```

```
plt.figure(figsize=(12, 8))
plt.imshow(image_np)
plt.axis('off')
plt.show()

# Test the function
detect_objects_in_image("path/to/test/image.jpg")
# Example placeholder for evaluation
test_loss = 0
num_batches = 0

for batch in val_dataset:
    predictions = detection_model(batch['image'], training=False)
    test_loss += detection_model.compute_loss(predictions, batch['groundtruth_boxes'],
batch['groundtruth_classes'])
    num_batches += 1

test_loss /= num_batches
print(f"Test Loss: {test_loss.numpy()}")
```